

GANs

Adán González Rodríguez
Sara Porto Álvarez

Código:

[Link al código](#)

Dataset

Explicación del Dataset:

Atributo	Valor
Número de clases	10
Tamaño de imagen	28 x 28 píxeles (grises)
Número de imágenes de entrenamiento	60,000
Número de imágenes de prueba	10,000
Formato	Escala de grises (1 canal)
Tipo de datos	Imagen + Etiqueta (label)

Dataset:

Time is 0.014961719512939453 sec

Shape of loading one batch: torch.Size([128, 1, 28, 28])

Total no. of batches present in trainloader: 469

Training Images



Generar

+ Código

+ Mar

GAN

Red neuronal del Generator:

```
[...]
```

Layer (type)	Output Shape	Param #
ConvTranspose2d-1	[-1, 256, 3, 3]	230,656
BatchNorm2d-2	[-1, 256, 3, 3]	512
ReLU-3	[-1, 256, 3, 3]	0
ConvTranspose2d-4	[-1, 128, 6, 6]	524,416
BatchNorm2d-5	[-1, 128, 6, 6]	256
ReLU-6	[-1, 128, 6, 6]	0
ConvTranspose2d-7	[-1, 64, 13, 13]	73,792
BatchNorm2d-8	[-1, 64, 13, 13]	128
ReLU-9	[-1, 64, 13, 13]	0
ConvTranspose2d-10	[-1, 1, 28, 28]	1,025
Tanh-11	[-1, 1, 28, 28]	0

=====
Total params: 830,785
Trainable params: 830,785
Non-trainable params: 0

Input size (MB): 0.00
Forward/backward pass size (MB): 0.42
Params size (MB): 3.17
Estimated Total Size (MB): 3.59

Generator(

Red neuronal del Crítico:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 13, 13]	1,088
BatchNorm2d-2	[-1, 64, 13, 13]	128
LeakyReLU-3	[-1, 64, 13, 13]	0
Conv2d-4	[-1, 128, 5, 5]	131,200
BatchNorm2d-5	[-1, 128, 5, 5]	256
LeakyReLU-6	[-1, 128, 5, 5]	0
Conv2d-7	[-1, 1, 1, 1]	2,049

Total params: 134,721

Trainable params: 134,721

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.32

Params size (MB): 0.51

Estimated Total Size (MB): 0.84

Critic(
 (disc): Sequential(
 (0): Sequential(
 (0): Conv2d(1, 64, kernel_size=(4, 4), stride=(2, 2))
 (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 ...

Optimizador:

```
lr = 0.0002
beta_1 = 0.5
beta_2 = 0.999

def weights_init(m):
    if isinstance(m, nn.Conv2d) or isinstance(m, nn.ConvTranspose2d):
        torch.nn.init.normal_(m.weight, 0.0, 0.02)
    if isinstance(m, nn.BatchNorm2d):
        torch.nn.init.normal_(m.weight, 0.0, 0.02)
        torch.nn.init.constant_(m.bias, 0)

gen = Generator(z_dim).to(device)
gen_opt = torch.optim.Adam(gen.parameters(), lr=lr, betas=(beta_1, beta_2))

crit = Critic().to(device)
crit_opt = torch.optim.Adam(crit.parameters(), lr=lr, betas=(beta_1, beta_2))

gen = gen.apply(weights_init)
crit = crit.apply(weights_init)
```

Gradient penalty:

```
def gradient_penalty(gradient):  
    gradient = gradient.view(len(gradient), -1)  
  
    gradient_norm = gradient.norm(2, dim=1)  
  
    penalty = torch.mean((gradient_norm - 1)**2)  
    return penalty
```

Loss:

```
def get_gen_loss(crit_fake_pred):  
    gen_loss = -1. * torch.mean(crit_fake_pred)  
    return gen_loss
```


```
def get_crit_loss(crit_fake_pred, crit_real_pred, gp, c_lambda):  
    crit_loss = torch.mean(crit_fake_pred) - torch.mean(crit_real_pred) + c_lambda * gp  
    return crit_loss
```

Resultados


Resultados de una ejecución:

Epoch 1,	Loss: 99.22,	Accuracy: 77.46%
Epoch 2,	Loss: 62.71,	Accuracy: 86.03%
Epoch 3,	Loss: 49.16,	Accuracy: 88.72%
Epoch 4,	Loss: 40.58,	Accuracy: 90.78%
Epoch 5,	Loss: 33.01,	Accuracy: 92.31%
Epoch 6,	Loss: 26.89,	Accuracy: 93.97%
Epoch 7,	Loss: 20.41,	Accuracy: 95.06%
Epoch 8,	Loss: 16.66,	Accuracy: 96.03%
Epoch 9,	Loss: 12.76,	Accuracy: 96.98%
Epoch 10,	Loss: 9.83,	Accuracy: 97.91%
Epoch 11,	Loss: 7.72,	Accuracy: 98.32%
Epoch 12,	Loss: 4.40,	Accuracy: 99.22%
Epoch 13,	Loss: 2.76,	Accuracy: 99.51%

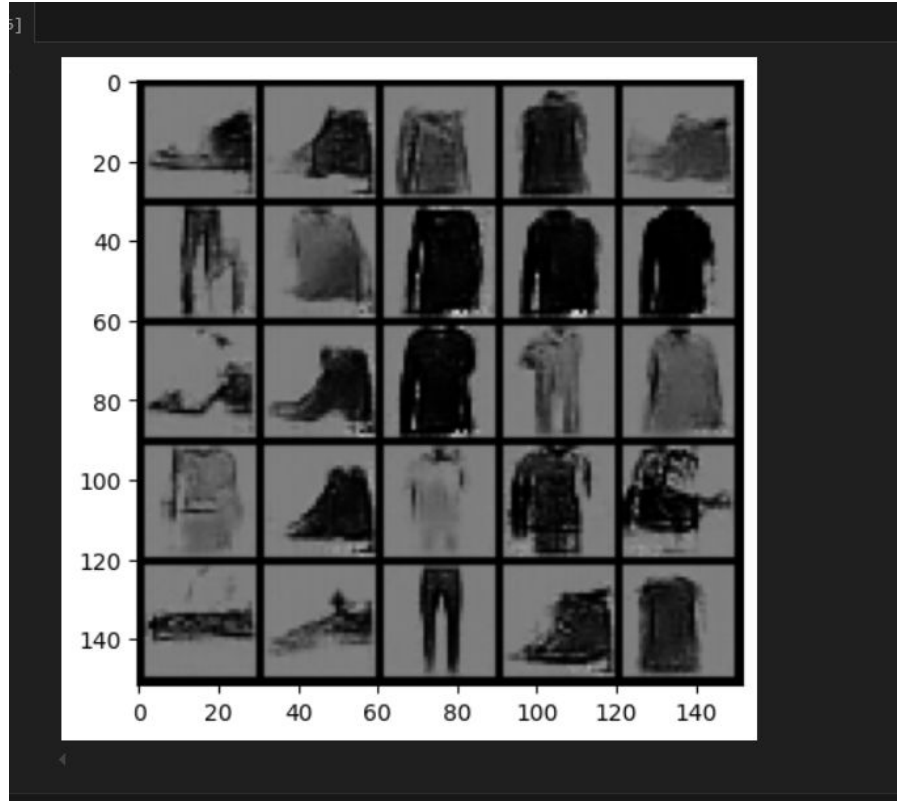
Training Loss (Generator and Critic) Batches:



Training Loss (Generator and Critic) Epochs:



Ejemplo de Generación de Imágenes con la GAN:



Entrenamiento del Clasificador

Clasificador:


```
class Classifier(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.conv = nn.Sequential(  
            nn.Conv2d(1, 32, 3, 1), nn.ReLU(),  
            nn.Conv2d(32, 64, 3, 1), nn.ReLU(),  
            nn.MaxPool2d(2)  
        )  
        self.fc = nn.Sequential(  
            nn.Flatten(),  
            nn.Linear(9216, 128), nn.ReLU(),  
            nn.Linear(128, 10)  
        )  
  
    def forward(self, x):  
        x = self.conv(x)  
        return self.fc(x)
```

Entrenamiento del clasificador:


Epoch 1,	Loss: 99.22,	Accuracy: 77.46%
Epoch 2,	Loss: 62.71,	Accuracy: 86.03%
Epoch 3,	Loss: 49.16,	Accuracy: 88.72%
Epoch 4,	Loss: 40.58,	Accuracy: 90.78%
Epoch 5,	Loss: 33.01,	Accuracy: 92.31%
Epoch 6,	Loss: 26.89,	Accuracy: 93.97%
Epoch 7,	Loss: 20.41,	Accuracy: 95.06%
Epoch 8,	Loss: 16.66,	Accuracy: 96.03%
Epoch 9,	Loss: 12.76,	Accuracy: 96.98%
Epoch 10,	Loss: 9.83,	Accuracy: 97.91%
Epoch 11,	Loss: 7.72,	Accuracy: 98.32%
Epoch 12,	Loss: 4.40,	Accuracy: 99.22%
Epoch 13,	Loss: 2.76,	Accuracy: 99.51%

Distribución por clases del resultado del clasificador:

- 0 – T-shirt/top
- 1 – Trouser
- 2 – Pullover
- 3 – Dress
- 4 – Coat
- 5 – Sandal
- 6 – Shirt
- 7 – Sneaker
- 8 – Bag
- 9 – Ankle boot



Confianza Obtenida de la generación de imágenes:



Resultados adicionales:

[Link a resultados adicionales](#)

Fin