

GANs

Adán González Rodríguez
Sara Porto Álvarez

Sobre el código:


[Link al código](#)

Explicación del Dataset:

| Atributo | Valor |
|-------------------------------------|----------------------------|
| Número de clases | 10 |
| Tamaño de imagen | 28 x 28 píxeles (grises) |
| Número de imágenes de entrenamiento | 60,000 |
| Número de imágenes de prueba | 10,000 |
| Formato | Escala de grises (1 canal) |
| Tipo de datos | Imagen + Etiqueta (label) |

Distribución por clases del Dataset:

- 0 – T-shirt/top
- 1 – Trouser
- 2 – Pullover
- 3 – Dress
- 4 – Coat
- 5 – Sandal
- 6 – Shirt
- 7 – Sneaker
- 8 – Bag
- 9 – Ankle boot



Red neuronal del Generator:

```
[...]  
...  
-----  
Layer (type)           Output Shape           Param #  
-----  
ConvTranspose2d-1      [-1, 256, 3, 3]        230,656  
BatchNorm2d-2          [-1, 256, 3, 3]         512  
ReLU-3                 [-1, 256, 3, 3]         0  
ConvTranspose2d-4      [-1, 128, 6, 6]        524,416  
BatchNorm2d-5          [-1, 128, 6, 6]         256  
ReLU-6                 [-1, 128, 6, 6]         0  
ConvTranspose2d-7      [-1, 64, 13, 13]       73,792  
BatchNorm2d-8          [-1, 64, 13, 13]        128  
ReLU-9                 [-1, 64, 13, 13]         0  
ConvTranspose2d-10     [-1, 1, 28, 28]        1,025  
Tanh-11                [-1, 1, 28, 28]         0  
-----  
Total params: 830,785  
Trainable params: 830,785  
Non-trainable params: 0  
-----  
Input size (MB): 0.00  
Forward/backward pass size (MB): 0.42  
Params size (MB): 3.17  
Estimated Total Size (MB): 3.59  
-----  
Generator(  
[...]
```

Red neuronal del Critico:

| Layer (type) | Output Shape | Param # |
|---------------|------------------|---------|
| Conv2d-1 | [-1, 64, 13, 13] | 1,088 |
| BatchNorm2d-2 | [-1, 64, 13, 13] | 128 |
| LeakyReLU-3 | [-1, 64, 13, 13] | 0 |
| Conv2d-4 | [-1, 128, 5, 5] | 131,200 |
| BatchNorm2d-5 | [-1, 128, 5, 5] | 256 |
| LeakyReLU-6 | [-1, 128, 5, 5] | 0 |
| Conv2d-7 | [-1, 1, 1, 1] | 2,049 |

Total params: 134,721

Trainable params: 134,721

Non-trainable params: 0

Input size (MB): 0.00

Forward/backward pass size (MB): 0.32

Params size (MB): 0.51

Estimated Total Size (MB): 0.84

Critic(
 (disc): Sequential(
 (0): Sequential(
 (0): Conv2d(1, 64, kernel_size=(4, 4), stride=(2, 2))
 (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
 ...

Resultados


Resultados de una ejecución

| | | |
|-----------|--------------|------------------|
| Epoch 1, | Loss: 99.22, | Accuracy: 77.46% |
| Epoch 2, | Loss: 62.71, | Accuracy: 86.03% |
| Epoch 3, | Loss: 49.16, | Accuracy: 88.72% |
| Epoch 4, | Loss: 40.58, | Accuracy: 90.78% |
| Epoch 5, | Loss: 33.01, | Accuracy: 92.31% |
| Epoch 6, | Loss: 26.89, | Accuracy: 93.97% |
| Epoch 7, | Loss: 20.41, | Accuracy: 95.06% |
| Epoch 8, | Loss: 16.66, | Accuracy: 96.03% |
| Epoch 9, | Loss: 12.76, | Accuracy: 96.98% |
| Epoch 10, | Loss: 9.83, | Accuracy: 97.91% |
| Epoch 11, | Loss: 7.72, | Accuracy: 98.32% |
| Epoch 12, | Loss: 4.40, | Accuracy: 99.22% |
| Epoch 13, | Loss: 2.76, | Accuracy: 99.51% |

Time is 0.014961719512939453 sec

Shape of loading one batch: torch.Size([128, 1, 28, 28])

Total no. of batches present in trainloader: 469




Generar


+ Código

+ Mar

Training Loss (Generator and Critic) Epochs




Training Loss (Generator and Critic) Batches



Ejemplo de Generación de Imágenes con la GAN



Confianza Obtenida de la generación de imágenes



Resultados adicionales:

[Link a resultados adicionales](#)